


# Deep Learning in Computer Vision



Alexei Manso Corrêa Machado

Pontifical Catholic University of Minas Gerais – D. Computer Science

# Unsupervised Learning

Steps:

- Choose a training set. The class of each element of the set is unknown
- Choose a set of features and the number of classes
- Choose a method for grouping elements
- Determine the parameters from the discovered classes
- Test with objects outside the training set.

# Supervised Learning

Steps:

- Choose a training set. The class of each element of the set is known.
- Choose discriminating characteristics.
- Choose a decision method / function.
- Determine the parameters from the training set.
- Test with objects outside the training set.

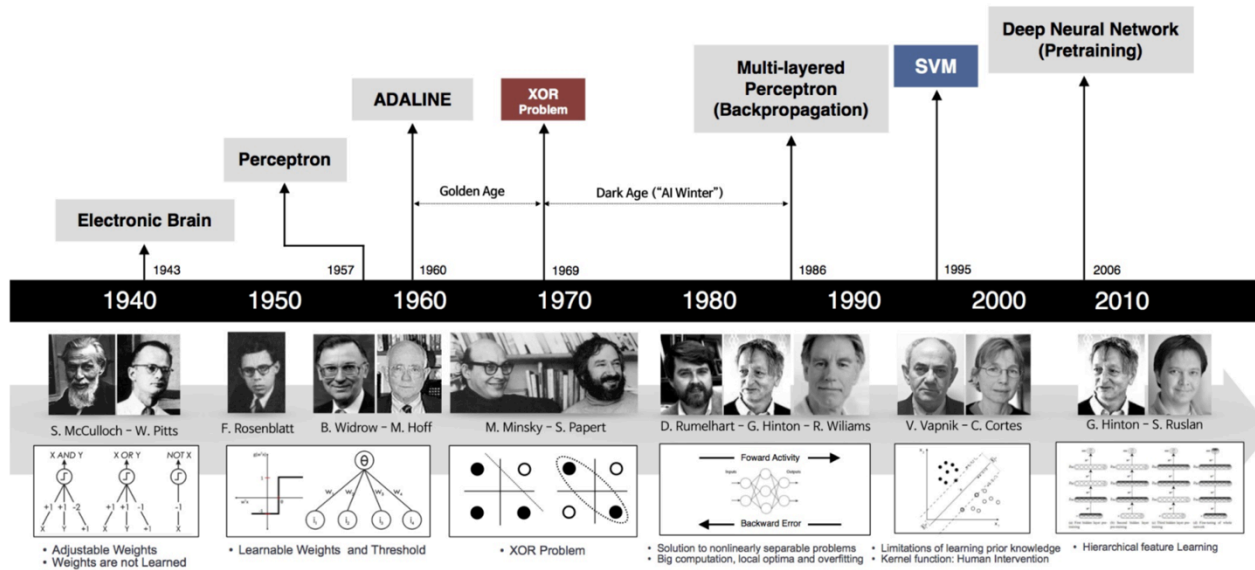
# Supervised Learning

Bayes Rule, the optimum decision function:

$$P(\mathcal{Y} | x_1 \dots x_n) = P(x_1 \dots x_n | \mathcal{Y})P(\mathcal{Y}) / P(x_1 \dots x_n)$$

- Minimum classification error
- The PDFs are however hard to determine in most cases
- Machine learning algorithms such as SVM, decision trees and neural networks try to approximate Bayes rule



# Milestones



# Feature selection


- Classical AI and Machine Learning algorithms requires the selection of features from an infinite set
- This is a problem of exponential complexity
- Heuristic solutions imply deeper knowledge of the problem domain
- Features should be uncorrelated
- Too many features can be a confounding aspect
- The “curse of dimensionality”
- Many classifiers require that the number of features be less than the number of training data
- DL algorithms are less sensitive to many of these issues

# The Perceptron (Rosenblatt 1958)




$$z = w_0 + w_1 x_1 + w_2 x_2 + \dots + w_n x_n = b + \mathbf{w}^T \mathbf{x}$$

$$a = \hat{y} = \sigma(z)$$



# The Perceptron (Rosenblatt 1958)

- The perceptron is a binary classifier:  $y=0$  if  $z>0$  otherwise  $y=1$
- The perceptron defines a linear decision function



- Example with 2 variables: the OR operator
- $b = -1$  (bias) 
$$z = -1 + 2x_1 + 3x_2$$



# The Perceptron (Rosenblatt 1958)

- Solutions also can be found for the AND and NOT operator
- What about XOR?

# The Perceptron (Rosenblatt 1958)

- How do we find a feasible solution?

1. Initialize  $W$

2. Repeat until  $W$  is stable (convergence)

- 2.1. For each sample  $(x,y)$  in the dataset

- 2.1.1 Compute  $\hat{y} = \sigma(b + w^T x)$

- 2.1.2. If  $\hat{y} \neq y$  then adjust each weight  $w_i$  so that  $\hat{y}$  gets closer to  $y$

- I.e. We backpropagate the output error in order to get a better estimate of the weights
- If the classes are linearly separable, the algorithm will converge, otherwise it will not!

# The Perceptron (Rosenblatt 1958)

## Remarks:

- A Sum of Squared error function can be used to evaluate convergence (Be careful to overfitting!)
- If the classes are linearly separable, the algorithm will converge, otherwise it will not! (Duda, Hart e Stork)
- The algorithm finds ANY solution that makes it converge. The SVM is an evolution of the perceptron that finds a decision function with maximum separability (Krauth e Mezard, 1987)
- If some input variable is useless its weight should have a small magnitude
- If we increase the number of input variables, linear separability may be achieved (not always, of course!)


# The Perceptron (Rosenblatt 1958)

Remarks:

- In 1969 Minsky and Pappert showed the weakness of the Perceptron to solve the XOR problem
- The research on Connectionism slowed down
- The perceptron model then receives 2 modifications: a different activation function and additional layers.

# Logistic regression

- In many cases, the behavior of a variable does not change drastically from one class to the other




- Therefore we want the output of the classifier to give the probability of the class instead of being a 0/1
- The probability however is not properly defined by a linear function!

# Logistic regression

We model the log of the odds ratio as a linear function, from which we get the sigmoid activation function:

$$\ln \frac{P(y=1)}{1-P(y=1)} = b + \mathbf{w}^T \mathbf{x}$$



# Logistic regression

- How to adapt the perceptron's learning algorithm to the new activation function?
  - We need to properly update the weights
  - We need to compute a loss function and know when to stop iterating
- For the logistic activation function, the SSE loss function does not work well
- Since  $P(y=1|x)=\hat{y}$  and  $P(y=0|x)=1-\hat{y}$ , we minimize the *cross-entropy loss function*:

$$L(y, \hat{y}) = -\ln P(y | x) = -\ln(\hat{y}^y (1 - \hat{y})^{1-y}) = -(y \ln \hat{y} + (1 - y) \ln(1 - \hat{y}))$$

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m L(y^{(i)}, \hat{y}^{(i)})$$

# A Bit of Calculus

Derivatives:

- Slope of the tangent of a function at a given value
- Variation rate of a function



# A Bit of Calculus

Chain rule:

■ Derivative of a composite function:

$$\text{If } F(x)=f(g(x)), F'(x)=f'(g(x))g'(x)$$

If we call  $u=g(x)$ ,  $dF/dx = dF/du * du/dx$

# A Bit of Calculus

Partial derivative:

- Derivative of a multivariate function with respect to each of the variables separately, the others being considered constant


$$\frac{\partial}{\partial x_i} f(x_1, \dots, x_n)$$

# A Bit of Calculus

Gradient:

- Vector of partial derivatives that indicates the direction of maximum variation


$$\nabla f(x_1, \dots, x_n) = [\partial f(x_1, \dots, x_n) / \partial x_1, \dots, \partial f(x_1, \dots, x_n) / \partial x_n]^T$$



En.wikipedia.org

# Computational graphs

- A model to compute the partial derivatives of the loss function with respect to the input, following the chain rule in backward direction
- Used to implement the backpropagation algorithm in some frameworks
- Ex:  $J=3(a+bc)$



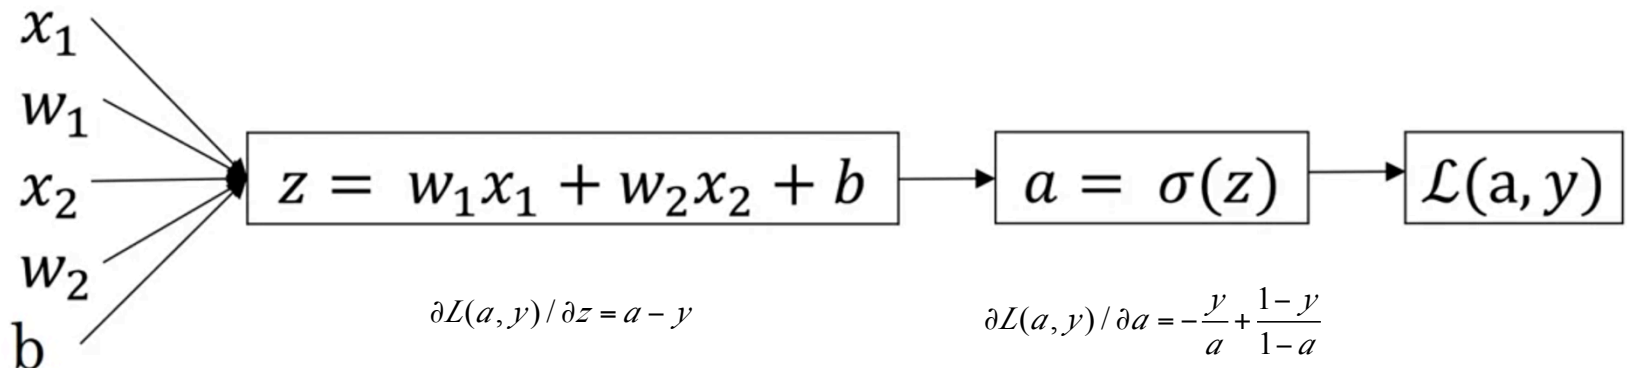
# Computational graphs

- Computational graph for logistic Loss function:

$$z = w^T x + b$$

$$\hat{y} = a = \sigma(z)$$

$$\mathcal{L}(a, y) = -(y \log(a) + (1 - y) \log(1 - a))$$



$$\partial L(a, y) / \partial w_i = x_i \partial L(a, y) / \partial z$$

$$\partial L(a, y) / \partial b = \partial L(a, y) / \partial z$$

# Gradient descent

1. Initialize parameters / Define hyperparameters
2. Loop for num\_iterations:
  - a. Forward propagation
  - b. Compute cost function
  - c. Backward propagation
  - d. Update parameters (using parameters, and grads from backprop)
4. Use trained parameters to predict labels

# Gradient descent

```
J = 0, dw1 = 0, dw2 = 0, db = 0
for i = 1 to m:
    z(i) = wTx(i) + b
    a(i) = σ(z(i))
    J += -[y(i) log a(i) + (1 - y(i)) log(1 - a(i))]
    dz(i) = a(i) - y(i)
    dw1 += x1(i) dz(i)
    dw2 += x2(i) dz(i)
    db += dz(i)
J = J/m, dw1 = dw1/m, dw2 = dw2/m
db = db/m
w1 -= (αdw1); w2 -= (αdw2); b -= (αdb)
```

Complexity?

# References and acknowledgements

Some of these slides were inspired or adapted from courses and presentations given by Andrew Ng, Camila Laranjeira, Fei-Fei Li, Flávio Figueiredo, Hugo Oliveira, Jefersson dos Santos, Justin Johnson, Keiller Nogueira, Pedro Olmo, Renato Assunção, Serena Yeung.

Reference courses include *Machine Learning* and *Deep Learning* CS230 and CS231 from Stanford University, *Deep Learning* and *Hands-on Deep Learning* from UFMG, *Deep Learning* CS498 from Un. Of Illinois.